

Product Search Relevance

Rohan Sadale, Akshay Kulkarni,
Gautham Sunder and Utkarsh
Kajaria



Data

- 240k (query,product) tuples
- 24k unique queries
- 97k unique products
- Product info:
 - Title
 - Description
 - Attributes
 - Relevance:
 - A continuous number between 1 and 3.
 - 1 - lowest, 3 - highest

Search Term	Product Title	Product Description	Product Attributes	Relevance

Problem Description

Given a search term, can we predict the relevance score of product?

Approach

- Classification
 - Two-class
 - Three-class
- Regression

Text Preprocessing

- Replace symbols and patterns
- Remove/replace words in parentheses with dots/space
- Unit Converter
 - pounds | pound | lbs | lb | lb. => "lb"
 - wattage | watts | watt => "watt."
- Remove commas between digits
 - 10, 000 will be replaced with 10000.
- Number to Digit Mapper
 - "one" => 1; "two" => 2.

```
def str_stem(s):  
    s = s.lower()  
    s = s.replace(",", "")  
    s = re.sub('(?[<=[0-9])[\ ]*centimeter[s]*(?=\ |$|\.)', '-cm ', s)  
    s = s.replace("vynal", "vinyl")  
    s = stemmer.stem(s)  
    return s
```

Feature Engineering

- Generated 26 features
- A combination of
 - basic descriptive features
 - distance features
 - word2vec embedding features
 - intersect count
 - position features
- Distance between search term and product title/product description
 - Jaccard Distance
 - Edit Distance

Feature Engineering

- Intersection based features
 - First and Last
 - Intersect count
 - Intersect position
- Word2Vec and Doc2Vec
 - Average Similarity
 - Centroid RMSE
- Longest Match
- TF-IDF Scores
- Asymmetric features that are non contextual such as presence of product dimensions

Models

- Linear
 - NaiveBayes
 - Logistic Regression
- Ensemble
 - Adaboost
 - Random Forests
- Neural Networks
- Regression

Experimental Setup


- Parameter Tuning
 - Logistic Regression $\lambda = [0.001, 0.01, 0.1, 1, 10, 1000]$
 - Random Forests
 - num_estimators [300, 500]
 - criterion : ['gini', 'entropy']
 - Max_features : ['auto', 'log2', None]
 - AdaBoost
 - criterion : ['gini', 'entropy']
 - max_features : ['auto', 'log2', None]

Experimental Setup

- Parameter Tuning
 - Neural Network
 - Num_hidden_layers : [1,2,3]
 - Activation : ['relu','sigmoid']
 - Regression - XGBoostRegressor
 - n_estimators : [100,150,200]
 - Max_depth : [3,5,7,9]
- Model Evaluation
 - Cross Validation
 - F - measure


Results (Two-class classification)

Binary Classification



Results (Three-class classification)

Ternary Classification



Results (Regression)

XgBoost Regression yielded an RMSE 0.47636 on the test data

Finished in the top 25% of Kaggle leaderboard (~2100 participants)



Thank you

Questions?

